

Found 11230 images belonging to 6 classes.

Found 2804 images belonging to 6 classes.

Found 3000 images belonging to 6 classes.

Explanation:

ImageDataGenerator: This is used to load images from the directory and apply preprocessing (in this case, just rescaling).

flow_from_directory: Loads images from the directory and assigns labels based on folder names.

We visualize 9 random images from the training set using matplotlib.


```
In [7]: import matplotlib.pyplot as plt
import numpy as np

class_names = list(train_data.class_indices.keys())


# Retrieve a batch of images and labels from the training data
images, labels = next(train_data)

# Plot some of the images
plt.figure(figsize=(10, 10))
for i in range(9):
    plt.subplot(3, 3, i + 1)
    plt.imshow(images[i])
    plt.title(class_names[np.argmax(labels[i])])
    plt.axis('off')
plt.show()
```


forest



buildings



glacier




sea




street




buildings



buildings



buildings



street



STEP 2: Model Architecture

Define a CNN architecture with convolutional layers, max pooling, batch normalization, and dropout.

```
In [8]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dr

        # Define the CNN model
        model = Sequential([
            Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
            MaxPooling2D((2, 2)),
            BatchNormalization(),

            Conv2D(64, (3, 3), activation='relu'),
            MaxPooling2D((2, 2)),
            BatchNormalization(),
```

```
Conv2D(128, (3, 3), activation='relu'),
MaxPooling2D((2, 2)),
BatchNormalization(),

Flatten(),
Dense(256, activation='relu'),
Dropout(0.5),
Dense(6, activation='softmax') # 6 classes in the dataset
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a

# Display the model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
batch_normalization (Batch Normalization)	(None, 74, 74, 32)	128
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 36, 36, 64)	256
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
batch_normalization_2 (Batch Normalization)	(None, 17, 17, 128)	512
flatten (Flatten)	(None, 36992)	0
dense (Dense)	(None, 256)	9470208
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 6)	1542
Total params: 9565894 (36.49 MB)		
Trainable params: 9565446 (36.49 MB)		
Non-trainable params: 448 (1.75 KB)		

Explanation:

We use Conv2D layers with ReLU activation, followed by MaxPooling2D to reduce spatial dimensions.

BatchNormalization helps in stabilizing and speeding up training.

The fully connected (dense) layers at the end allow the model to make predictions based on the learned features.

We use softmax activation in the final layer since this is a multi-class classification problem.

STEP 3: Model Training

Train the model on the training data and validate it on the validation data.

```
In [9]: # Train the model with training and validation data
history = model.fit(
    train_data,
    validation_data=validation_data,
    steps_per_epoch=train_data.samples // 32,
    validation_steps=validation_data.samples // 32,
    epochs=20
)
```

Epoch 1/20
350/350 [=====] - 96s 270ms/step - loss: 5.4759 - accuracy: 0.4294 - val_loss: 2.4049 - val_accuracy: 0.3825
Epoch 2/20
350/350 [=====] - 94s 268ms/step - loss: 1.5079 - accuracy: 0.5143 - val_loss: 1.2872 - val_accuracy: 0.5269
Epoch 3/20
350/350 [=====] - 94s 268ms/step - loss: 1.1509 - accuracy: 0.5699 - val_loss: 0.9456 - val_accuracy: 0.6498
Epoch 4/20
350/350 [=====] - 93s 266ms/step - loss: 1.0941 - accuracy: 0.5932 - val_loss: 0.9199 - val_accuracy: 0.6494
Epoch 5/20
350/350 [=====] - 93s 267ms/step - loss: 1.0641 - accuracy: 0.6050 - val_loss: 0.9405 - val_accuracy: 0.6724
Epoch 6/20
350/350 [=====] - 93s 266ms/step - loss: 1.0376 - accuracy: 0.6204 - val_loss: 1.0035 - val_accuracy: 0.6455
Epoch 7/20
350/350 [=====] - 93s 265ms/step - loss: 1.0051 - accuracy: 0.6234 - val_loss: 1.0657 - val_accuracy: 0.6200
Epoch 8/20
350/350 [=====] - 94s 269ms/step - loss: 0.9785 - accuracy: 0.6407 - val_loss: 0.8951 - val_accuracy: 0.6642
Epoch 9/20
350/350 [=====] - 94s 268ms/step - loss: 0.9578 - accuracy: 0.6498 - val_loss: 0.8715 - val_accuracy: 0.6828
Epoch 10/20
350/350 [=====] - 93s 266ms/step - loss: 0.9325 - accuracy: 0.6557 - val_loss: 1.1449 - val_accuracy: 0.5858
Epoch 11/20
350/350 [=====] - 94s 268ms/step - loss: 0.9285 - accuracy: 0.6631 - val_loss: 1.0149 - val_accuracy: 0.6422
Epoch 12/20
350/350 [=====] - 94s 269ms/step - loss: 0.9125 - accuracy: 0.6738 - val_loss: 0.8896 - val_accuracy: 0.6800
Epoch 13/20
350/350 [=====] - 93s 265ms/step - loss: 0.8736 - accuracy: 0.6836 - val_loss: 0.8427 - val_accuracy: 0.6997
Epoch 14/20
350/350 [=====] - 94s 269ms/step - loss: 0.8664 - accuracy: 0.6831 - val_loss: 0.8889 - val_accuracy: 0.6731
Epoch 15/20
350/350 [=====] - 94s 267ms/step - loss: 0.8296 - accuracy: 0.7018 - val_loss: 1.1103 - val_accuracy: 0.6494
Epoch 16/20
350/350 [=====] - 93s 265ms/step - loss: 0.8144 - accuracy: 0.7161 - val_loss: 0.8949 - val_accuracy: 0.6778
Epoch 17/20
350/350 [=====] - 94s 269ms/step - loss: 0.8035 - accuracy: 0.7178 - val_loss: 0.9964 - val_accuracy: 0.6239
Epoch 18/20
350/350 [=====] - 94s 268ms/step - loss: 0.7787 - accuracy: 0.7264 - val_loss: 1.1395 - val_accuracy: 0.6096
Epoch 19/20
350/350 [=====] - 94s 268ms/step - loss: 0.7712 - a

```
ccuracy: 0.7324 - val_loss: 0.7498 - val_accuracy: 0.7446
Epoch 20/20
350/350 [=====] - 93s 267ms/step - loss: 0.7387 - a
ccuracy: 0.7402 - val_loss: 0.9594 - val_accuracy: 0.6789
```

model.fit() trains the CNN using the training.

STEP 4: Model Evaluation


Evaluate the model on the test set and plot training and validation accuracy/loss.

```
In [10]: # Evaluate on test data
test_loss, test_acc = model.evaluate(test_data)
print(f"Test accuracy: {test_acc}")

# Plot accuracy and loss curves
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.title('Accuracy')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Loss')
plt.show()
```

```
94/94 [=====] - 4s 43ms/step - loss: 1.0381 - accur
acy: 0.6960
Test accuracy: 0.6959999799728394
```



We evaluate the trained model on the test data to get the test accuracy.

The accuracy and loss curves help visualize the model's performance across epochs.

STEP 5: Display Confusion Matrix

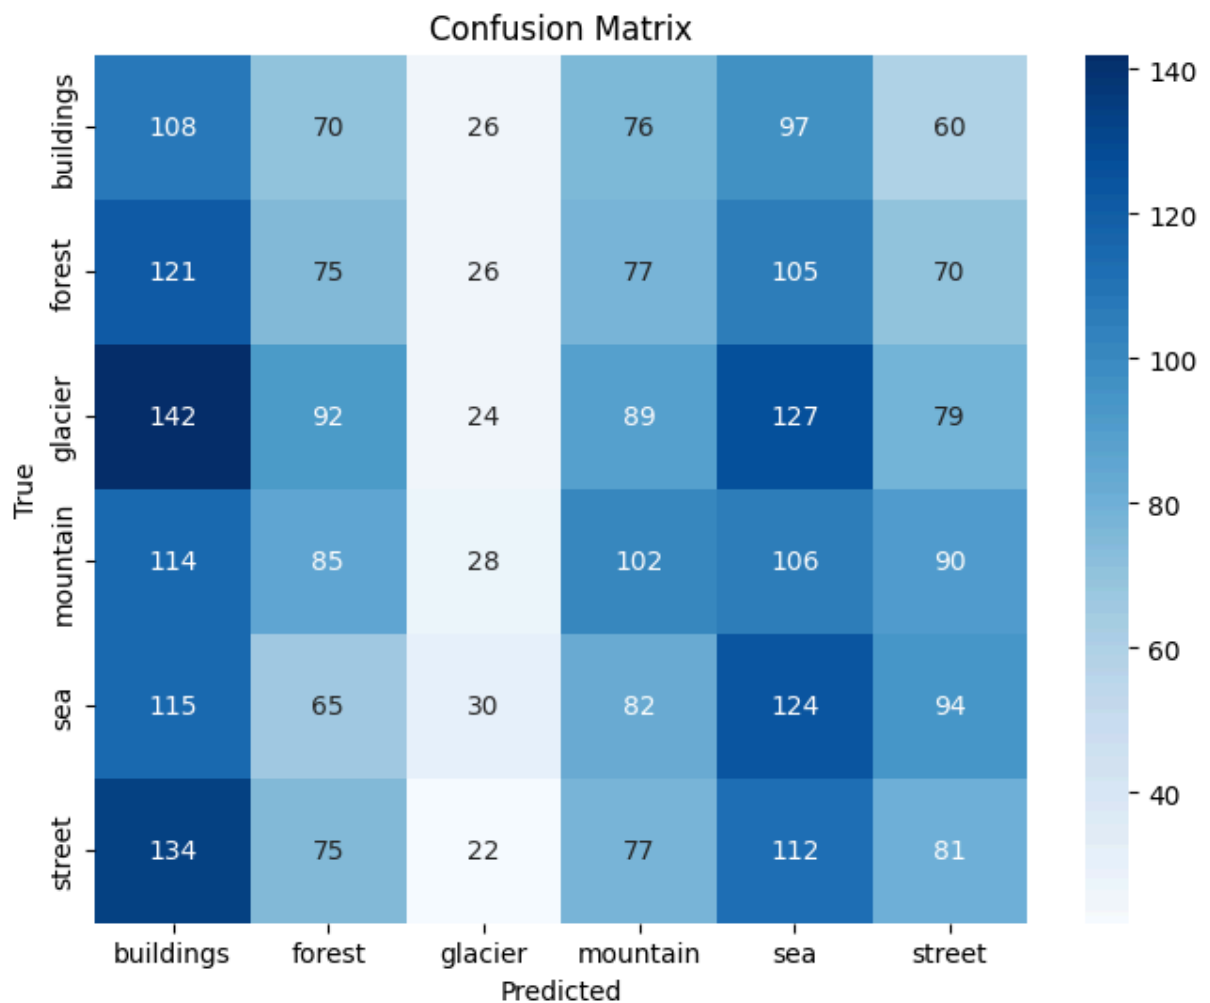
Analyze the model's predictions on the test set with a confusion matrix.


```
In [11]: from sklearn.metrics import confusion_matrix
import seaborn as sns

# Get predictions on test data
y_pred = np.argmax(model.predict(test_data), axis=-1)
y_true = test_data.classes

# Plot confusion matrix
cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=class_names,
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix")
plt.show()
```

94/94 [=====] - 4s 44ms/step



STEP 5: Optimization

1. Experiment with Data Augmentation
2. Fine-Tune Hyperparameters
 - 2.1. Adjust Learning Rate and Batch Size
 - 2.2. Modify the Number of Filters

```
In [ ]: # Define ImageDataGenerator with additional augmentation techniques
train_datagen = ImageDataGenerator(
    rescale=1./255,          # Rescale pixel values
    validation_split=0.2,    # Split for validation set
    rotation_range=40,       # Random rotations
    width_shift_range=0.2,   # Randomly shift images horizontally
    height_shift_range=0.2,  # Randomly shift images vertically
    shear_range=0.2,        # Shear transformations
    zoom_range=0.2,         # Random zoom
    horizontal_flip=True,    # Randomly flip images
    fill_mode='nearest'     # Fill mode for new pixels
)

# Load training and validation data with the updated augmentation
train_data = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

validation_data = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
)
```

```
In [ ]: from tensorflow.keras.optimizers import Adam

# Compile the model with a different learning rate
learning_rate = 1e-4 # Reduced learning rate
model.compile(optimizer=Adam(learning_rate=learning_rate),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Re-train the model with a different batch size
history = model.fit(
    train_data,
    validation_data=validation_data,
    steps_per_epoch=train_data.samples // 32, # Adjust this according to the
    validation_steps=validation_data.samples // 32,
    epochs=20
)
```

```

In [ ]: # Define a new CNN model with different filter sizes
model = Sequential([
    Conv2D(64, (3, 3), activation='relu', input_shape=(150, 150, 3)), # Inc
    MaxPooling2D((2, 2)),
    BatchNormalization(),

    Conv2D(128, (3, 3), activation='relu'), # Increased filters
    MaxPooling2D((2, 2)),
    BatchNormalization(),

    Conv2D(256, (3, 3), activation='relu'), # Increased filters
    MaxPooling2D((2, 2)),
    BatchNormalization(),

    Flatten(),
    Dense(512, activation='relu'), # Increased dense layer size
    Dropout(0.5),
    Dense(6, activation='softmax') # Output layer for 6 classes
])

# Compile the new model
model.compile(optimizer=Adam(learning_rate=learning_rate),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Re-train the model
history = model.fit(
    train_data,
    validation_data=validation_data,
    steps_per_epoch=train_data.samples // 32,
    validation_steps=validation_data.samples // 32,
    epochs=20
)

```